

Lecture 3. Vectors

R and Data Visualization

BIG2006, Hanyang University, Fall 2022

Scalars, Vectors, Arrays, and Matrices

R variables types are called *modes*.

All elements in a vector must have the same mode, which can be integer, numeric (floating-point number), character (string), logical (Boolean), complex, and so on.

Note: If you need your program code to check the mode of a variable `x`, you can query it by the call `typeof(x)`.

Adding and Deleting Vector Elements

- ▶ Vectors are stored like arrays in C, so you cannot insert or delete element.
- ▶ The size of vector is determined at its creation.
⇒ To add or delete elements, reassign the vector!

```
x <- c(1983, 5, 15, 2022)
```

```
x <- c(x[1:3], 9999, x[4])
```

```
x
```

```
## [1] 1983    5    15 9999 2022
```

Obtaining the Length of a Vector, `length()`

```
x <- c(1,2,10,5000)
length(x)
```

```
## [1] 4
```

- ▶ We often need to use `length()`.
- ▶ Now we want to have a function that determines the index of the first "1" value in the function's vector argument (assuming we are sure there is such a value).
- ▶ How do you write the code?

```
first1 <- function(x){  
  for (i in 1:length(x)){  
    if (x[i]==1) break # break out of loop  
  }  
  return(i)  
}
```

Note: Without the `length()` function, we would have needed to add a second argument to `first1()`, say `n`, to specify the length of `x`.

Matrices and Arrays as Vectors

- Arrays and matrices are actually vectors too.

```
m <- matrix(1:4,nrow=2,ncol=2,byrow=TRUE)
```

```
m
```

```
##      [,1] [,2]
```

```
## [1,]    1    2
```

```
## [2,]    3    4
```

```
m + 10:13 # (1,3,2,4) + (10,11,12,13) = (11,14,14,17)
```

```
##      [,1] [,2]
```

```
## [1,]   11   14
```

```
## [2,]   14   17
```

Declarations

Compiled languages requires that you *declare* variables.

C example:

```
int x;  
int y[3];
```

In R, you do not declare variables.

```
z <- 3  
y <- vector(length=2)  
y[1] <- 5  
y[2] <- 12 # or y <- c(5,12)
```

Recycling

When applying an operation to two vectors that requires them to be the same length, R automatically recycles or repeats, the shorter one.

```
> c(6,0,9,20,22) + c(1,2,4)
```

```
[1] 7 2 13 21 24
```

```
Warning message:
```

```
In c(6, 0, 9, 20, 22) + c(1, 2, 4) :
```

```
  longer object length is not a multiple of shorter  
object length
```


Common Vector Operations

We will cover arithmetic and logical operations, vector indexing, and some useful ways to create vectors.

Vector Arithmetic and Logical Operations

- ▶ Scalars are actually one-elements vectors.
- ▶ $+$, $-$, $*$, $/$, and $\%\%$ operations will be applied element-wise.

```
2+3
```

```
## [1] 5
```

```
"+"(2,3)
```

```
## [1] 5
```

```
x <- c(1,5,9)
```

```
x + c(5,-5,0)
```

```
## [1] 6 0 9
```

```
x * c(5,0,-1)
```

```
## [1] 5 0 -9
```

```
x / c(5,4,-1)
```

```
## [1] 0.20 1.25 -9.00
```

```
x %% c(5,4,-1)
```

```
## [1] 1 1 0
```

Vector Indexing

- Form a subvector by picking elements of the given vector for specific indices

```
y <- c(1.2, 3.9, 0.4, 0.12)
y[c(1,3)] # extract elements 1 and 3 of y
## [1] 1.2 0.4
```

```
v <- 2:3
y[v]
## [1] 3.9 0.4
```

- Negative subscript: exclude the given elements in the output

```
z <- c(5,112,500)
z[-1] # exclude element 1
```

```
## [1] 112 500
```

```
z[1:length(z)-1] # 1:(length(z)-1) ?
```

```
## [1] 5 112
```

```
z[-length(z)]
```

```
## [1] 5 112
```

Generating Useful Vectors with the : Operator

- ▶ The colon operator : produces a vector consisting of a range of numbers.

```
c(5:8, "||", 5:1)
```

```
## [1] "5" "6" "7" "8" "||" "5" "4" "3" "2" "1"
```

```
i <- 3
```

```
c(1:i-1, "||", 1:(i-1))
```

```
## [1] "0" "1" "2" "||" "1" "2"
```

Generating Vector Sequences with seq()

- ▶ seq() is a generalization of : and generates a sequence in arithmetic progression.

```
seq(from=-10,to=10,by=2)
```

```
## [1] -10 -8 -6 -4 -2 0 2 4 6 8 10
```

```
seq(from=0.1,to=2,length=5)
```

```
## [1] 0.100 0.575 1.050 1.525 2.000
```

```
x <- c(9,15,2022)
seq(x)
```

```
## [1] 1 2 3
```

```
x <- NULL
x
```

```
## NULL
```

```
seq(x)
```

```
## integer(0)
```

Note: `seq(x)` gives us the same results as `1 : length(x)` if `x` is not empty, but it correctly evaluates to `NULL` if `x` is empty.

Repeating Vector Constants with rep()

- ▶ `rep()` allows us to put the same constant into long vectors.

```
x <- rep(7,5)
```

```
x
```

```
## [1] 7 7 7 7 7
```

```
rep(c(5,15,2019),3)
```

```
## [1] 5 15 2019 5 15 2019 5 15 2019
```

```
rep(c(3:1,2))
```

```
## [1] 3 2 1 2
```

Using `all()` and `any()`

The `any()` and `all()` functions are handy shortcuts.

They report whether any or all of their arguments are `TRUE`.

```
x <- 1:15  
c(any(x>8), any(x>20))
```

```
## [1] TRUE FALSE
```

```
c(all(x>8), all(x>0))
```

```
## [1] FALSE TRUE
```

Vectorized Operations

Suppose we have a function $f()$ that we wish to apply to all elements of a vector x .

In many cases, we can accomplish this by simply calling $f()$ on x itself.

This can really simplify our code and give us a dramatic performance increase of hundredfold or more.

One of the most effective ways to achieve speed in R code is to use operations that are *vectorized*, meaning that a function applied to a vector is actually applied individually to each element.

Vector In, Vector Out

- ▶ Vectorized operations ($+$, $*$, $>$) enable a potential speedup.

```
u <- c(5,2,8)
```

```
v <- c(-1,3,9)
```

```
u > v
```

```
## [1] TRUE FALSE FALSE
```

```
w <- function(x) return(x+1)
```

```
w(u)
```

```
## [1] 6 3 9
```

- ▶ Even the transcendental functions, square roots, logs, trig functions, and so on, are vectorized.

```
sqrt(1:5)
```

```
## [1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

```
y <- c(1.2, 3.9, 0.4)
```

```
z <- round(y)
```

```
z
```

```
## [1] 1 4 0
```

```
f <- function(x,c) return((x+c)^2)  
f(1:3,0)
```

```
## [1] 1 4 9
```

```
f(1:3,2)
```

```
## [1] 9 16 25
```

Vector In, Matrix Out

- ▶ Consider the function itself is vector-valued, as `z12()` is here:

```
z12 <- function(z) return(c(z,z^2))  
x <- 1:4  
z12(x)  
## [1] 1 2 3 4 1 4 9 16
```

⇒ More natural to have these arranged as an 9-by-2 matrix, which we can do with the `matrix` function.


```
matrix(z12(x),ncol=2)
```

```
##      [,1] [,2]  
## [1,]    1    1  
## [2,]    2    4  
## [3,]    3    9  
## [4,]    4   16
```

- ▶ We can streamline things using `sapply()` (simplify apply).
- ▶ The call `sapply(x,f)` applies the function `f()` to each element of `x` and then converts the result to a matrix.

```
z12 <- function(z) return(c(z,z^2))  
sapply(1:9,z12)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]  
## [1,]    1    2    3    4    5    6    7    8    9  
## [2,]    1    4    9   16   25   36   49   64   81
```

NA and NULL Values

In statistical data sets, we often encounter missing data, which we represent in R with the value NA.

NULL, on the other hand, represents that the value in question simply doesn't exist, rather than being existent but unknown.

Using NA

```
x <- c(15, NA, 99, 2019)
```

```
x
```

```
## [1] 15 NA 99 2019
```

```
c(mean(x), mean(x, na.rm=T)) # instruction to skip NAs
```

```
## [1] NA 711
```

```
x <- c(15, NULL, 99, 2019)
```

```
mean(x)
```

```
## [1] 711
```

Using NULL

built up a vector of the even numbers in 1:20

```
z <- NULL
```

```
for (i in 1:20) if (i %% 2 == 0) z <- c(z,i)
```

```
z
```

```
## [1] 2 4 6 8 10 12 14 16 18 20
```

```
z <- NA
```

```
for (i in 1:20) if (i %% 2 == 0) z <- c(z,i)
```

```
z
```

```
## [1] NA 2 4 6 8 10 12 14 16 18 20
```

Note: NULL values really are counted as nonexistent and NULL is a special R object with no mode.

Filtering

Another feature reflecting the functional programming nature of R is *filtering*.

This allows us to extract a vector's elements that satisfy certain conditions.

Generating Filtering Indices

► Example 1

```
z <- c(5, 2, -3, 8)
w <- z[z*z > 8]
w
## [1] 5 -3 8
```

Note: Evaluation of the expression $z * z > 8$ gives us a vector of Boolean values!

► Example 2

```
z <- c(5, 2, -3, 8)
y <- c(1, 2, 30, 5)
y[z*z > 8]
```

```
## [1] 1 30 5
```


► Example 3 (involving assignment)

```
x <- c(1, 3, 8, 2, 20)
x[x > 3] <- 0
x
```

```
## [1] 1 3 0 2 0
```

Filtering with the subset() Function

- ▶ The difference between subset() and ordinary filtering lies in the manner in which NA values are handled.

```
x <- c(6, 1:3, NA, 12)
```

```
x
```

```
## [1]  6  1  2  3 NA 12
```

```
x[x > 5]
```

```
## [1]  6 NA 12
```

```
subset(x, x > 5)
```

```
## [1]  6 12
```

The Selection Function `which()`

- ▶ `which()` can extract the positions of elements in the vector at which the condition occurs.

```
z <- c(5, 2, -3, 8)
which(z*z < 8)

## [1] 2
```

Note: $z * z < 8$ is evaluated to (FALSE, TRUE, FALSE, FALSE).

A Vectorized if-then-else: The ifelse() Function

```
x <- 1:10  
y <- ifelse(x %% 2 == 0, 5, 12) # %% is the mod operator  
y
```

```
## [1] 12  5 12  5 12  5 12  5 12  5
```

Here, we wish to produce a vector in which there is a 5 wherever x is even or a 12 wherever x is odd.

```
x <- c(5,2,9,12)  
y <- ifelse(x > 6, 2*x, 3*x) # %% is the mod operator  
y
```

```
## [1] 15  6 18 24
```

Testing Vector Equality

Suppose we want to test whether two vectors are equal.

```
x <- 1:3  
y <- c(1,3,4)  
x == y
```

```
## [1] TRUE FALSE FALSE
```

Note: In fact, `==` is a vectorized function. The expression `x == y` applies the function `== ()` to the elements of x and y , yielding a vector of Boolean values.

```
"=="(3,2)
```

```
## [1] FALSE
```

```
i <- 2
```

```
"=="(i,2)
```

```
## [1] TRUE
```

```
all(x == y)
```

```
## [1] FALSE
```

```
identical(x,y)
```

```
## [1] FALSE
```

Vector Element Names

The elements of a vector can optionally be given names.

```
x <- c(5,15,2019)
names(x)
```

```
## NULL
```

```
names(x) <- c("day", "month", "year")
names(x)
```

```
## [1] "day"    "month"  "year"
```

```
x
```

```
##    day month  year
##     5    15  2019
```

```
x["month"]
```

```
## month
```

```
##      15
```

```
names(x) <- NULL
```

```
x
```

```
## [1]      5      15 2019
```


More on c()

If the arguments you pass to `c()` are of differing modes, they will be reduced to a type that is the lowest common denominator, as follows:

```
c(5,2,"hanyang")
```

```
## [1] "5"      "2"      "hanyang"
```

R consider the list mode to be of lower precedence in mixed expressions.

```
c(5,2,list(math=3,stat=7))
```

```
## [[1]]
```

```
## [1] 5
```

```
##
```

```
## [[2]]
```

```
## [1] 2
```

```
##
```

```
## $math
```

```
## [1] 3
```

```
##
```

```
## $stat
```

```
## [1] 7
```

Another point to keep in mind is that `c()` has a flattening effect for vectors.

```
c(5,2,c(1.8,3.5))
```

```
## [1] 5.0 2.0 1.8 3.5
```

Reference

- ▶ Matloff, N. [The Art of R Programming: A Tour of Statistical Software Design](#). No Starch Press. Chapter 2.